

# Storing XML on Relational Tables

Federico Ulliana  
GraphIK, LIRMM, INRIA

Slides collected from James Cheney and Sam Idicula

Boston, winter '99 : XML standardization

Jan 2000 : people wondering ...

*Now, how can I publish online my relational data?*  
(XMLAGG - Xperanto; Mappings - SilkRoute)

Boston, winter '99 : XML standardization

Jan 2000 : people wondering ...

*Now, how can I publish online my relational data?  
(XMLAGG - Xperanto; Mappings - SilkRoute)*

Feb 2000 : people (again) wondering ...

*I created my first 10GB XML document crawling web data.  
Now, how can I query it ?*

# 3 schools for processing XML data

1. Flat streams: store XML data as is in text files
  - query support: limited; fast for retrieving whole documents
2. Native XML Databases: designed specifically for XML
  - XML document stored in XML specific way
  - Goal: Efficient support for XML queries
3. Re-use existing DB storage systems
  - Leverage mature systems (DBMS)
  - How ? Map XML document into flat tables

# Why transform XML data into relations?

Native XML databases need:

- storing XML data, indexing,
- query processing/optimization
- concurrency control
- updates
- access control, . . .
- **Nontrivial**: the study of these issues is still in its infancy – incomplete support for general data management tasks

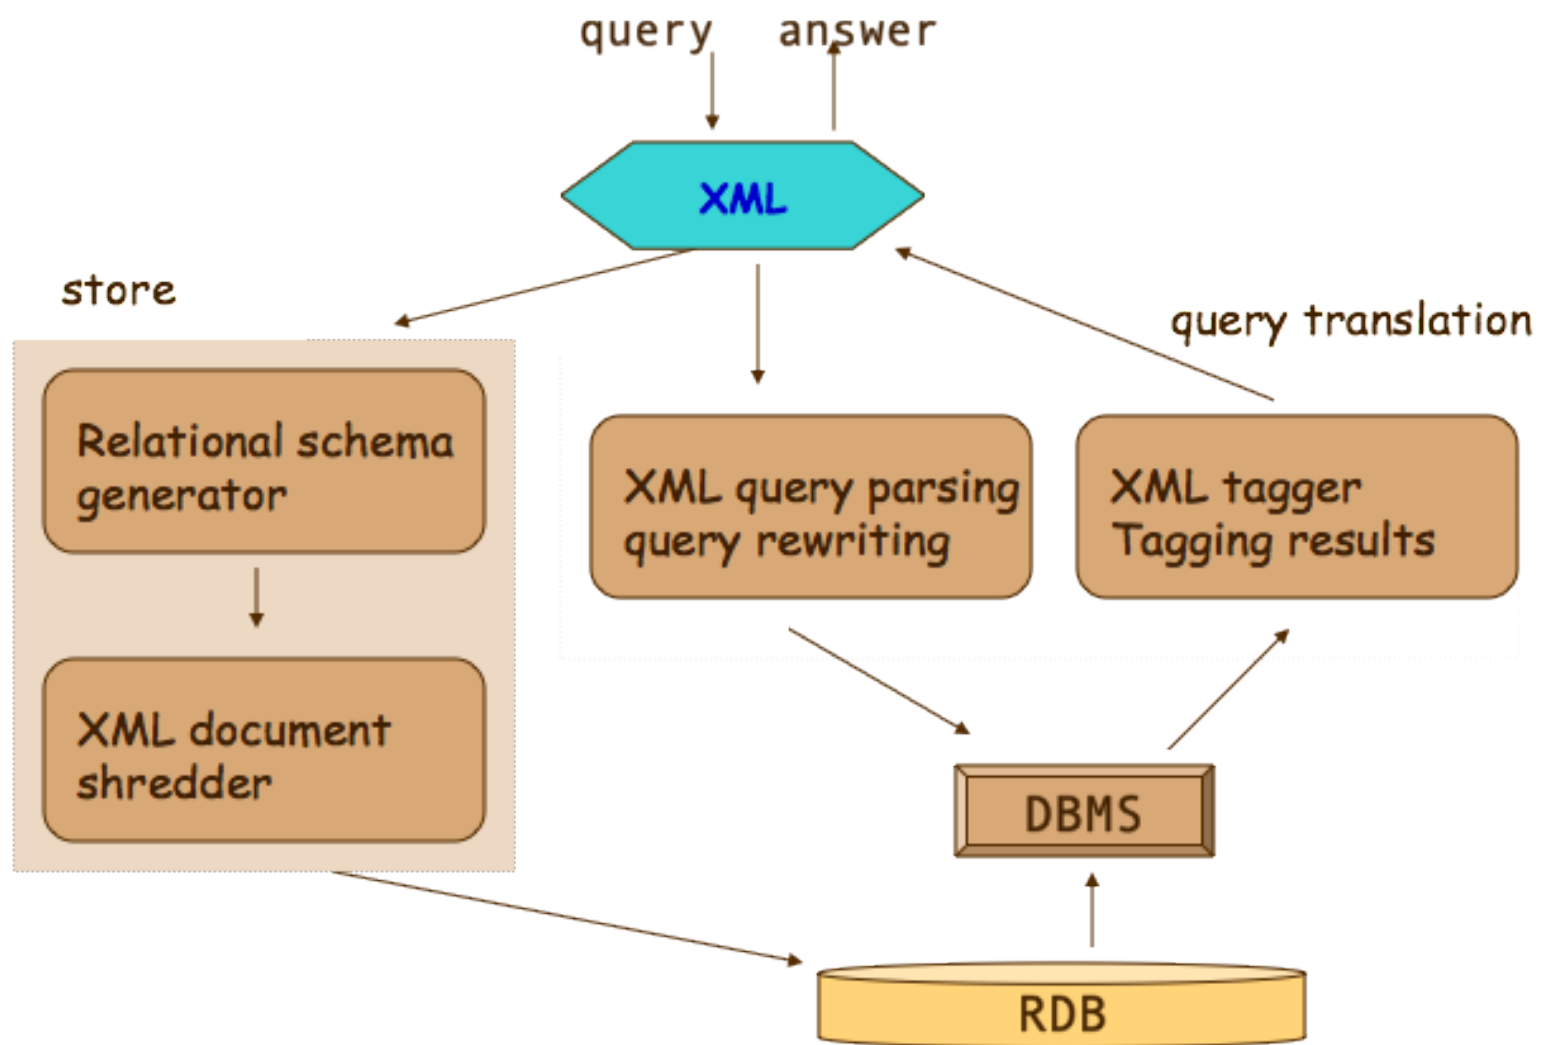
Haven't these already been developed for relational DBMS!?

- Why not take advantage of available DBMS techniques?

# From XML to relations : steps

1. Derive a relational schema
2. Insert XML data into relational tuples
3. Translate XML queries to SQL queries
4. Convert query results back to XML

# Architecture



# Nontrivial issues

## Data model mismatch

- DTD: recursive, regular expressions/nested content
- relational schema: tables, single-valued attributes

## Information preservation

- lossless: there should be an effective method to reconstruct the original XML document from its relational storage
- propagation/preservation of integrity constraints

## Query language mismatch

- XQuery, XSLT: Turing-complete
- XPath: transitive edges (descendant, ancestor)
- SQL: first-order, limited / no recursion



# Plan

Schema-unaware

Schema-aware

Commercial solutions

# **SCHEMA-UNAWARE XML STORAGE**

# Schema-unaware storage

Storage easier if we have a fixed schema

But, often don't have schema

Or schema may change over time

- schema updates require reorganizing or reloading!

So: schema-oblivious XML storage

**Schema chaos:** In this scenario, customers want the flexibility to manage XML data that may or may not have schema, or may have “any” schema. For instance, a telecommunication customer wants to manage XML data generated from different towers, which generate documents with slightly different schemas from each other. They want to store them in one table and perform efficient query on the shared common pieces.

# The basics first

*“before thinking about sophisticated solutions,  
how the simplest and most obvious approaches  
perform?”*

Round 1)      EDGE   vs   VERTICAL-EDGE

# EDGE storage


Observation: XML ordered trees can be encoded with

binary relation

EDGE (parent, child)

order relation

NEXT-SIBLING (prec, succ)




EDGE (  $n_1$ ,  $n_2$  )

EDGE (  $n_1$ ,  $n_3$  )

NEXT-SIBLING (  $n_2$ ,  $n_3$  )

# Edges & Values



## EDGES

source	target	ordinal	tag	type
	<b>n<sub>1</sub></b>		people	elt
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1	person	elt
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2	person	elt
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3	person	elt
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	name	elt
<b>n<sub>3</sub></b>	<b>n<sub>4</sub></b>	1		txt
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1	name	elt
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1	age	elt
<b>n<sub>8</sub></b>	<b>n<sub>9</sub></b>	1		num


## TEXTVALUES

node	value
<b>n<sub>4</sub></b>	Serge

## NUMVALUES

node	value
<b>n<sub>9</sub></b>	42


# Querying



Q = /people/person/age/text( )




# Querying



$Q = /people/person/age/text()$

/people/person/age/text()



## EDGES

source	target	ordinal	tag	type
	<b>n<sub>1</sub></b>		people	elt
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1	person	elt
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2	person	elt
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3	person	elt
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	name	elt
<b>n<sub>3</sub></b>	<b>n<sub>4</sub></b>	1		txt
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1	name	elt
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1	age	elt
<b>n<sub>8</sub></b>	<b>n<sub>9</sub></b>	1		num

## TEXTVALUES

node	value
<b>n<sub>4</sub></b>	Serge

## NUMVALUES

node	value
<b>n<sub>9</sub></b>	42

# /people/person/age/text( )

```
SELECT N.value
FROM    EDGES as e1
        EDGES as e2
        EDGES as e3
        EDGES as e4
        NUMVALUES N

WHERE

        e1.target=e2.source
AND      e2.target=e3.source
AND      e3.target=e4.source
AND      e1.tag="people"
AND      e2.tag="person"
AND      e3.tag="age"
AND      e3.target=e4.source
AND      e4.type="num"
AND      e4.target= N.node
```

## EDGES

source	target	ordinal	tag	type
	<b>n<sub>1</sub></b>		people	elt
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1	person	elt
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2	person	elt
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3	person	elt
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	name	elt
<b>n<sub>3</sub></b>	<b>n<sub>4</sub></b>	1		txt
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1	name	elt
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1	age	elt
<b>n<sub>8</sub></b>	<b>n<sub>9</sub></b>	1		num

## TEXTVALUES

node	value
<b>n<sub>4</sub></b>	Serge

## NUMVALUES

node	value
<b>n<sub>9</sub></b>	42

# /people/person/age/text ( )

```
SELECT N.value
FROM    EDGES as e1
        EDGES as e2
        EDGES as e3
        EDGES as e4
        NUMVALUES N
WHERE
        e1.target=e2.source
AND     e2.target=e3.source
AND     e3.target=e4.source
AND     e1.tag="people"
AND     e2.tag="person"
AND     e3.tag="age"
AND     e3.target=e4.source
AND     e4.type="num"
AND     e4.target= N.node
```

Lots of joins



TEXTVALUES

node	value
<b>n<sub>4</sub></b>	Serge

NUMVALUES

node	value
<b>n<sub>9</sub></b>	42

# /people/person/age/text()

```
SELECT N.value
FROM    EDGES as e1
        EDGES as e2
        EDGES as e3
        EDGES as e4
        NUMVALUES N
WHERE
        e1.target=e2.source
AND      e2.target=e3.source
AND      e3.target=e4.source
AND      e1.tag="people"
AND      e2.tag="person"
AND      e3.tag="age"
AND      e3.target=e4.source
AND      e4.type="num"
AND      e4.target= N.node
```

We also need a query  
testing for text values  
( UNION )

TEXTVALUES

node	value
<b>n<sub>4</sub></b>	Serge

NUMVALUES

node	value
<b>n<sub>9</sub></b>	42

# Querying

**Fragmentation:** tree  
spread across the table



EDGES

source	target	ordinal	tag	type
	<b>n<sub>1</sub></b>		people	
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1	person	ref
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2	person	ref
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3	person	ref
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	name	ref
<b>n<sub>3</sub></b>	<b>n<sub>4</sub></b>	1		txt
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1	name	ref
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1	age	ref
<b>n<sub>8</sub></b>	<b>n<sub>9</sub></b>	1		num

TEXTVALUES

node	value
<b>n<sub>4</sub></b>	Serge

NUMVALUES

node	value
<b>n<sub>9</sub></b>	42

# Querying

**Fragmentation:** tree spread across the table

Indexes **unaware** of tree structure

## EDGES

source	target	ordinal	tag	type
	<b>n<sub>1</sub></b>		people	
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1	person	ref
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2	person	ref
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3	person	ref
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	name	ref
<b>n<sub>3</sub></b>	<b>n<sub>4</sub></b>	1		txt
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1	name	ref
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1	age	ref
<b>n<sub>8</sub></b>	<b>n<sub>9</sub></b>	1		num

## TEXTVALUES

node	value
<b>n<sub>4</sub></b>	Serge

## NUMVALUES

node	value
<b>n<sub>9</sub></b>	42

# How to improve ?

## 1. Vertical partitioning

group edges targeting  
same tag-label

## 2. Inlining

put text and numeric  
values in the main table

EDGES

source	target	ordinal	tag	type
	<b>n<sub>1</sub></b>		people	
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1	person	ref
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2	person	ref
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3	person	ref
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	name	ref
<b>n<sub>3</sub></b>	<b>n<sub>4</sub></b>	1		txt
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1	name	ref
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1	age	ref
<b>n<sub>8</sub></b>	<b>n<sub>9</sub></b>	1		num

TEXTVALUES


node	value
<b>n<sub>4</sub></b>	Serge

NUMVALUES

node	value
<b>n<sub>9</sub></b>	42



# VERTICAL-EDGE + Inline



people

source	target	ordinal	txtval	numval
n <sub>1</sub>				

person

source	target	ordinal	txtval	numval
n <sub>1</sub>	n <sub>2</sub>	1		
n <sub>1</sub>	n <sub>5</sub>	2		
n <sub>1</sub>	n <sub>7</sub>	3		

name

source	target	ordinal	txtval	numval
n <sub>2</sub>	n <sub>3</sub>	1	Serge	
n <sub>5</sub>	n <sub>6</sub>	1		

age

source	target	ordinal	txtval	numval
n <sub>7</sub>	n <sub>8</sub>	1		42

# VERTICAL-EDGE + Inline

Q = /people/person/age/text()

```
SELECT AGE.value
FROM   PEOPLE  P1
       PERSON  P2
       AGE
```

WHERE

P1.target=P2.source

AND

P2.target=AGE.source

people

source	target	ordinal	txtval	numval
<b>n<sub>1</sub></b>				

person

source	target	ordinal	txtval	numval
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1		
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2		
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3		

name

source	target	ordinal	txtval	numval
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	Serge	
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1		

age


source	target	ordinal	txtval	numval
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1		42

Joins on smaller tables







# VERTICAL-EDGE+Inline beats EDGE

(query-answering time with the two storages)



# The queries SeFrWh you cannot ask

- Does it exists a direct flight between Paris and Los Angeles ? 
- Does it exists a (possibly indirect) flight between Montpellier and Austin ? 
  - problem : we do not know the number of intermediary airports (=joins)
- Does it exists a child for the node N ? 
- Is the node M a descendant of node N ? 
  - problem : we do not know the depth of a descendant node
  - taking max document depth is not an elegant solution

# Issues with XPath axes

Q = /people//age/text()

Descendant = implicit recursion  
sort of (child)\*

Does not translate to  
SELECT-FROM-WHERE query

Recursion :

ORACLE, POSTGRES      OK

MySQL                      NO

people

source	target	ordinal	txtval	numval
	<b>n<sub>1</sub></b>			

person

source	target	ordinal	txtval	numval
<b>n<sub>1</sub></b>	<b>n<sub>2</sub></b>	1		
<b>n<sub>1</sub></b>	<b>n<sub>5</sub></b>	2		
<b>n<sub>1</sub></b>	<b>n<sub>7</sub></b>	3		

name

source	target	ordinal	txtval	numval
<b>n<sub>2</sub></b>	<b>n<sub>3</sub></b>	1	Serge	
<b>n<sub>5</sub></b>	<b>n<sub>6</sub></b>	1		

age

source	target	ordinal	txtval	numval
<b>n<sub>7</sub></b>	<b>n<sub>8</sub></b>	1		42

# Limits of Edge/Vertical

Indexing unaware of tree structure

- fragmentation : subtree spread across db

Incomplete query translation


- descendant axis steps involve recursion

Lots of joins

- joins + no indexing = trouble

# MONET storage


(so called because developed first on Monet-DB)



**Idea** : one table for each path  
in the XML tree

- people
- people\_person
- people\_person\_name
- people\_person\_age

# MONET storage



people

node	txtval	numval
n <sub>1</sub>		

people\_person

node	txtval	numval
n <sub>2</sub>		
n <sub>5</sub>		
n <sub>7</sub>		

people\_person\_name

node	txtval	numval
n <sub>3</sub>	Serge	
n <sub>6</sub>		

people\_person\_age

node	txtval	numval
n <sub>8</sub>		42



/people/person/age/text()

```
SELECT txtval,numval  
FROM  
people_person_age
```

people

node	txtval	numval
<b>n<sub>1</sub></b>		

people\_person

node	txtval	numval
<b>n<sub>2</sub></b>		
<b>n<sub>5</sub></b>		
<b>n<sub>7</sub></b>		

people\_person\_name

node	txtval	numval
<b>n<sub>3</sub></b>	Serge	
<b>n<sub>6</sub></b>		

people\_person\_age

node	txtval	numval
<b>n<sub>8</sub></b>		42

# Performances

- MONET (obviously) beats VERTICAL-EDGE+Inlining

# Still one question...

And descendant axis ?

$Q = \text{/people//age}$

How to select the relations to query ?

/people//age

people\_(*any-seq*)\_age

# people\_(*any-seq*)\_age

people



people\_person



people\_person\_name



people\_person\_age



//person//\*

*(any-seq)\_person\_(any-seq)\_(any-tag)*



*(any-seq)\_person\_(any-seq)\_(any-tag)*

people



people\_person



people\_person\_name



people\_person\_age



*(any-seq)\_person\_(any-seq)\_(any-tag)*

```
SELECT node  
FROM people_person_name
```

**UNION**

```
SELECT node  
FROM people_person_age
```

people\_person\_name    people\_person\_age



**And the remaining axes ?**


Maybe we need some new ideas...



# **INTERVALS**

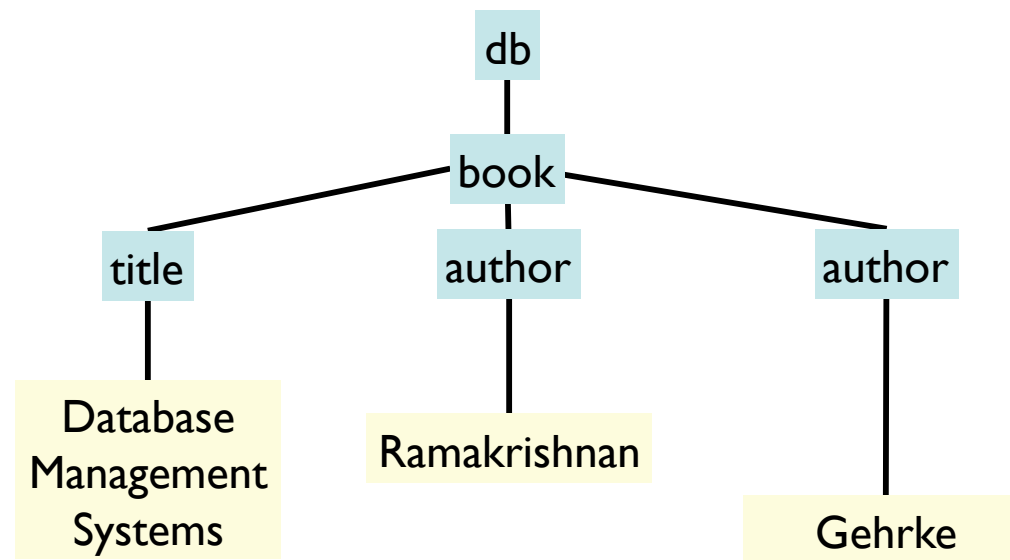
# Intervals

Idea: Node-identifier embed navigational-information



# Intervals

Think of XML as a linear string




```
<db><book><title>Database Management Systems</title><author>Ramakrishnan</author><author>Gehrke</author></book></db>
```

# Intervals

**Begin** : the first time we see a node (opening tag)


**End** : the last time we see a node (closing tag)



Each node corresponds to an interval on line





# Begin/end numbering



NODE Table

begin	end	par	tag	type
1	16		db	ELT
2	15	1	book	ELT
3	6	2	title	ELT
4	5	3		TEXT
7	10	2	author	ELT
8	9	7		TEXT
11	14	2	author	ELT
12	13	11		TEXT





End

15

10

5

ance + follow

desc+  
prec


0


5

10

15

Begin





# From Axes to Intervas


Child(**n**<sub>1</sub>, **n**<sub>2</sub>)

⇔

**n**<sub>1</sub>.begin = **n**<sub>2</sub>.par

NODE Table

begin	end	par	tag	type
<b>1</b>	16		db	ELT
2	15	<b>1</b>	book	ELT
3	6	2	title	ELT
4	5	3		TEXT
7	10	2	author	ELT
8	9	7		TEXT
11	14	2	author	ELT
12	13	11		TEXT



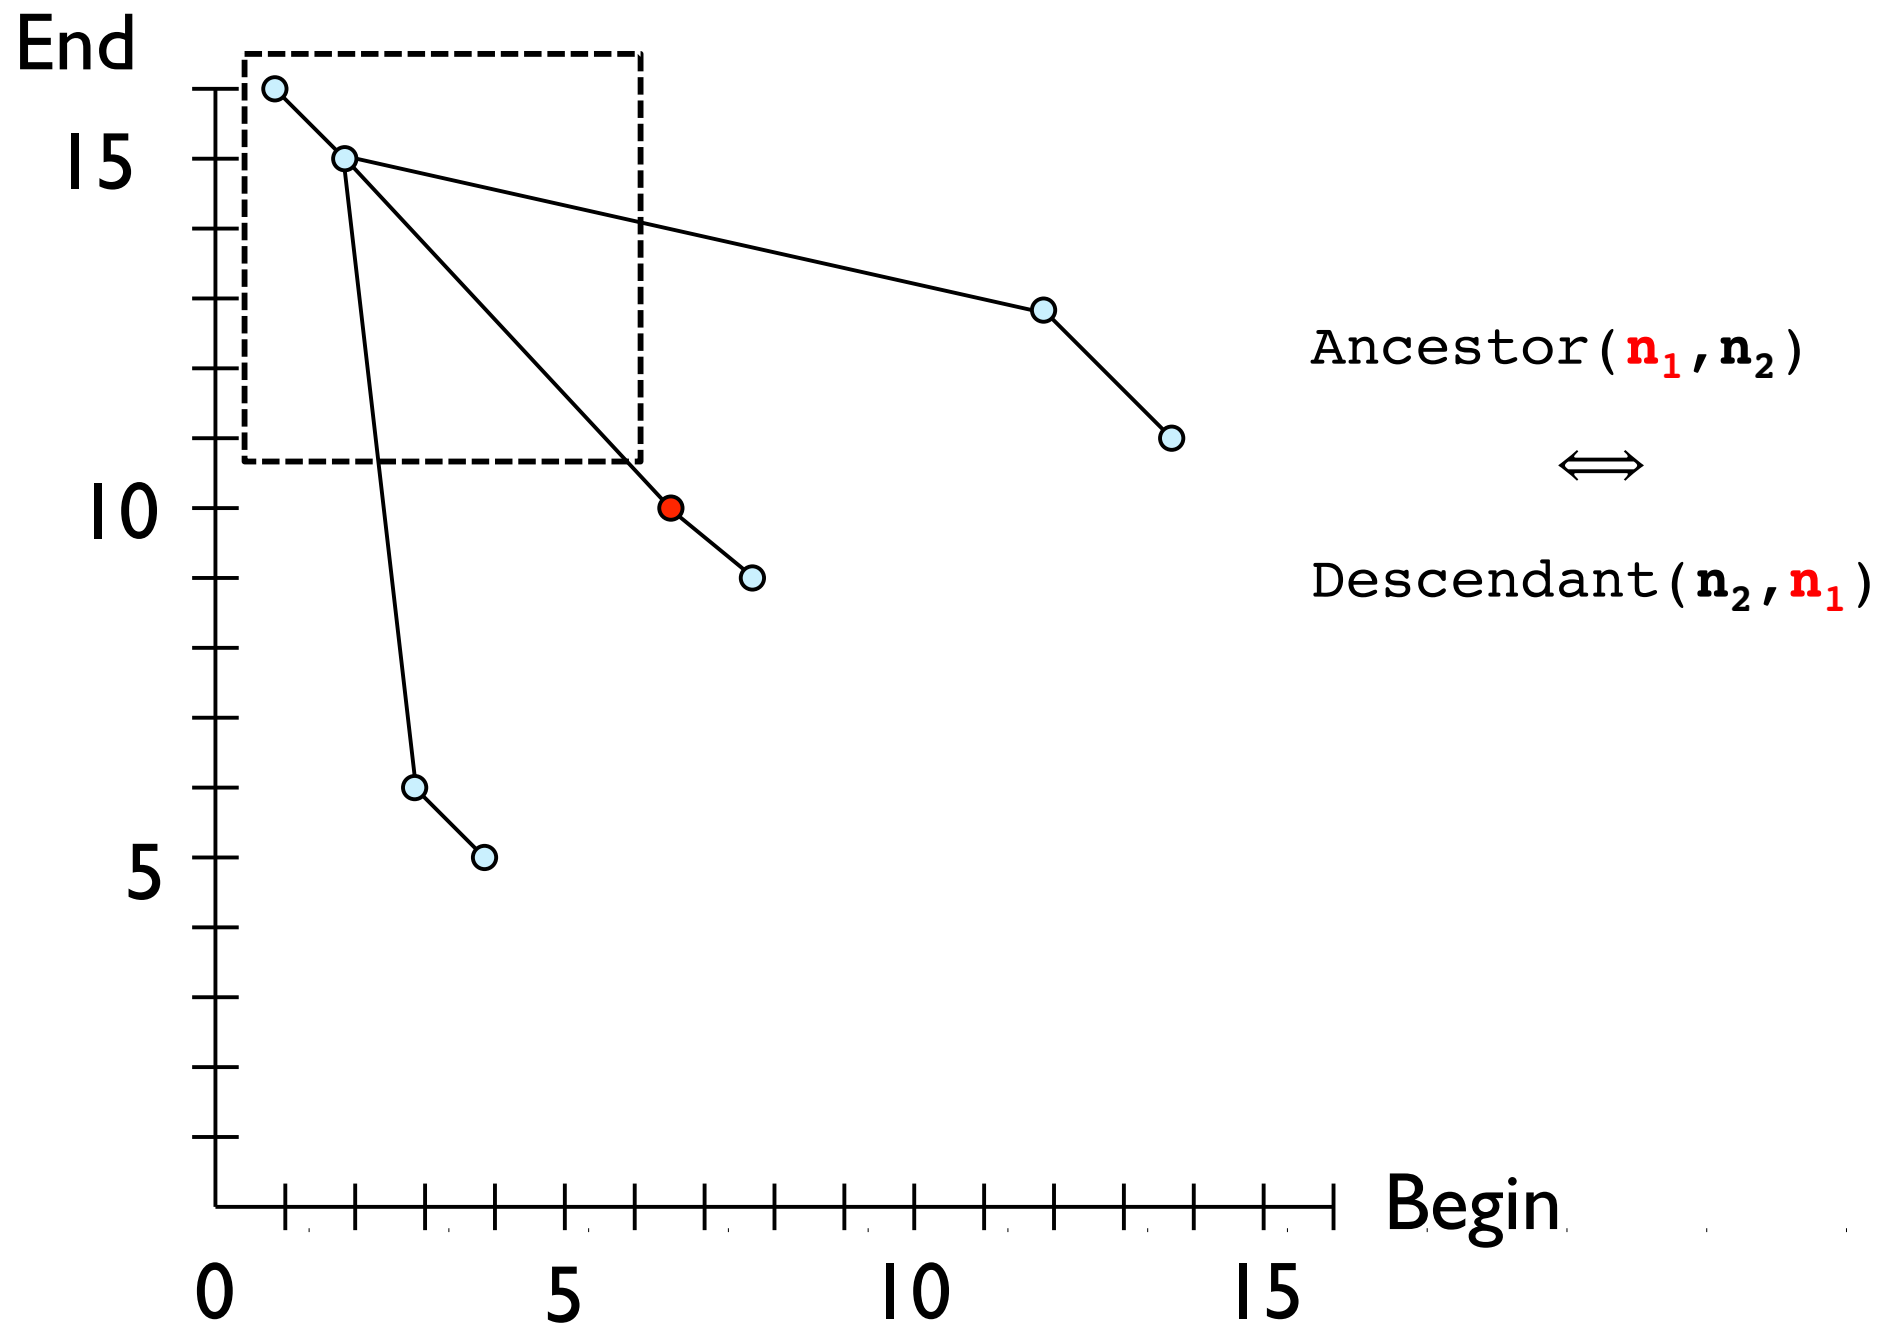
Descendant ( $\mathbf{n}_1, \mathbf{n}_2$ )

$\Leftrightarrow$


$\mathbf{n}_1.\text{begin} < \mathbf{n}_2.\text{begin}$


and

$\mathbf{n}_1.\text{end} > \mathbf{n}_2.\text{end}$









End

15

10

5


0

5

10

15

Begin



Prec-Sib(**n**<sub>1</sub>, **n**<sub>2</sub>)

⇔

**n**<sub>1</sub>.begin > **n**<sub>2</sub>.end

and

**n**<sub>1</sub>.par = **n**<sub>2</sub>.par

# Ready to Query (with all axes!)

```
Q = //a//b/ancestor::c//d/following-sibling::e
```

```
SELECT e.nodeID
```

```
FROM node a, node b, node c, node d, node e
```

```
WHERE
```

```
    a.tag = 'a', b.tag = 'b',  
    c.tag = 'c', d.tag='d', e.tag='e'
```

```
AND Descendant(a.nodeID,b.nodeID)
```

```
AND Ancestor(b.nodeId,c.nodeId)
```



```
AND Descendant(c.nodeId,d.nodeId)
```

```
AND Following-Sibling(d.nodeId,e.nodeId)
```

to simplify the query, we assume that the nodes have also a unique **nodeId**

# Other Approaches: Pre/Post

(Gurst et al. 2004)




Replaces Begin/End with Pre/Post visit of the tree

# Other Approaches: Dewey Decimal Encoding



Each node's ID is a list of integers

- $[i_1, i_2, \dots, i_n]$  (often written  $i_1.i_2. \dots .i_n$ )
- giving the "path" from root to this node



nodeID	tag	type
[]	db	ELT
1	book	ELT
1.1	title	ELT
1.1.1		TEXT
1.2	author	ELT
1.2.1		TEXT
1.3	author	ELT
1.3.1		TEXT

# Summary

Dewey: string index, requires PREFIX, LEN UDFs

Interval: integer begin/end, pre/post indexes, only requires arithmetic

What about updates?

- Dewey: requires renumbering (exist update-friendly variants)
- Interval encoding: can require re-indexing most of the document

# **SCHEMA-AWARE XML STORAGE**



# Derivation of relational schema from DTD

## Should be lossless

- the original document can be effectively reconstructed from its relational representation

## Should support querying

- XML queries should be able to be rewritten to efficient relational queries

# A book DTD

Complex  
Regular Expressions

```
<!ELEMENT db (book*)>
```

```
<!ELEMENT book (title,author*,chapter*, ref*)>
```

```
<!ELEMENT chapter (text | section)*>
```

```
<!ELEMENT ref book>
```

```
<!ELEMENT title #PCDATA>
```

```
<!ELEMENT author #PCDATA>
```

```
<!ELEMENT section #PCDATA>
```

```
<!ELEMENT text #PCDATA>
```

Recursion



# Recall :regular expressions

$r ::=$	$\epsilon$	empty sequence
	$ $	
	$a$	(tag element name)
	$ $	
	$(r, s)$	sequential composition
	$ $	
	$(r   s)$	union
	$ $	
	$(r^*)$	repetition

$$r^+ = r^*, r$$

$$r? = r | \epsilon$$

# First-step : Simplification of RegExp

$r ::= \epsilon$

$| a$


$| (r, s)$



Order does not matter

$| (r | s)$

$| (r1, r2)^*$



Correlation does not matter

$(a, b)^* \quad \text{--1--} \rightarrow \quad (a^*, b^*) \quad \text{--2--} \rightarrow \quad (a^* | b^*)$

# A book DTD

```
<!ELEMENT book (title,authors*,chapter*,ref*)>
```

```
<!ELEMENT chapter (text | section)*>
```

is approximated by

```
<!ELEMENT book (title|authors*|chapter*| ref*)>
```

```
<!ELEMENT chapter (text*) | (section*) >
```


still correct, but less precise

# Second step: create a graph representation of the DTD

```
<!ELEMENT book (title|author*|chapter*| ref*)>
```

```
<!ELEMENT chapter (text*) | (section*) >
```

```
<!ELEMENT ref book>
```




# Second step: create a graph representation of the DTD

```
<!ELEMENT book (title|author*|chapter*| ref*)>
```

```
<!ELEMENT chapter (text*) | (section*) >
```

```
<!ELEMENT ref book>
```




# Second step: create a graph representation of the DTD

```
<!ELEMENT book (title|author*|chapter*| ref*)>
```

```
<!ELEMENT chapter (text*) | (section*) >
```

```
<!ELEMENT ref book>
```





# Graph representation of DTD


1. Each element type / attribute is represented by a unique node
2. Edges represent the subelement (and attribute) relations
3. Symbol \*: denotes 0 or more occurrences of subelements
4. Cycles indicate recursion
  1. e.g., `book -> ref -> book -> ref`

# Third step: Create Relations + Inline

Traverse the DTD graph depth-first and create relations for

- (1) the root
- (2) each \* node
- (3) each recursive node
- (4) each node with at least 2 parents

Nodes (w/out \* and w/ only 1 parent) are inlined as fields: no relation created



# Third step: Create Relations + Inline

book(bookID, title: string)

author(authorID, author: string)


chapter(chapterID)

text(textID, text: string)

section(sectionID, section: string)

ref(refID)

we forgot  
something..



# Third step: Create Relations + Inline

book(bookID, title: string)


author(authorID, **bookID**, author: string)

chapter(chapterID, **bookID**)

text(textID, **chapterID**, text: string)

section(sectionID, **chapterID**, section: string)


ref(refID, **bookID**)



# Still missing detail : parent-ambiguity

```
book(bookID, code, title: string)  
author(authorID, bookID, author: string)  
chapter(chapterID, bookID)  
text(textID, chapterID)  
section(sectionID, chapterID)  
ref(refID, bookID)
```

code needed to distinguish book and  
ref parents



# Still missing detail : parent-ambiguity

book(bookID, **flag**, title: string)

author(authorID, bookID, author: string)

chapter(chapterID, bookID)

text(textID, chapterID, text: string)

section(sectionID, chapterID, section: string)

ref(refID, bookID)

Foreign keys:

book.parentID  $\subseteq$  db.dbID      if **flag**= 1

book.parentID  $\subseteq$  ref.refID    if **flag**= 0



text



section



ref

# Relational schema

`book(bookID, code, title: string)`

`author(authorID, bookID, author: string)`

`chapter(chapterID, bookID)`

`text(textID, chapterID, text: string)`

`section(sectionID, chapterID, section: string)`

`ref(refID, bookID)`

To preserve the semantics

- ID: each relation has an artificial ID (key)
- parentID: foreign key coding edge relation
- We can also add column naming path in the DTD graph

Note: `title` is inlined

# Summary of schema-ware XML

Use DTD/XML Schema to decompose document

Reorganization of regular expressions

- $(\text{text} , \text{section})^* \rightarrow \text{text}^* \mid \text{section}^*$
- document order and type-correlations are lost

Querying: Supports a large class of common XML queries

- Fast lookup & reconstruction of inlined elements
- Systematic translation unclear



**COMMERCIAL SOLUTIONS**

# Well, XML is just text, right?

Most databases allow CLOB (Character Large Object) columns - unbounded length string.

So you just store the XML text in one of these

Surprisingly popular

- and can make sense for storing "document-like" parts of XML data (eg HTML snippets)
- But not a good idea if you want to query the XML

# SQL / XML

Instead of blindly using CLOBs.. extend SQL with XML features

- "XML" column type
- XPath or XQuery queries (or updates) on XML columns

Also surprisingly popular (DB2, IBM, Oracle)

- Pro: At least DB knows it's XML
- Pro: Part of SQL 2003 (SQL/XML extensions)
- Con: Frankenstein's query language

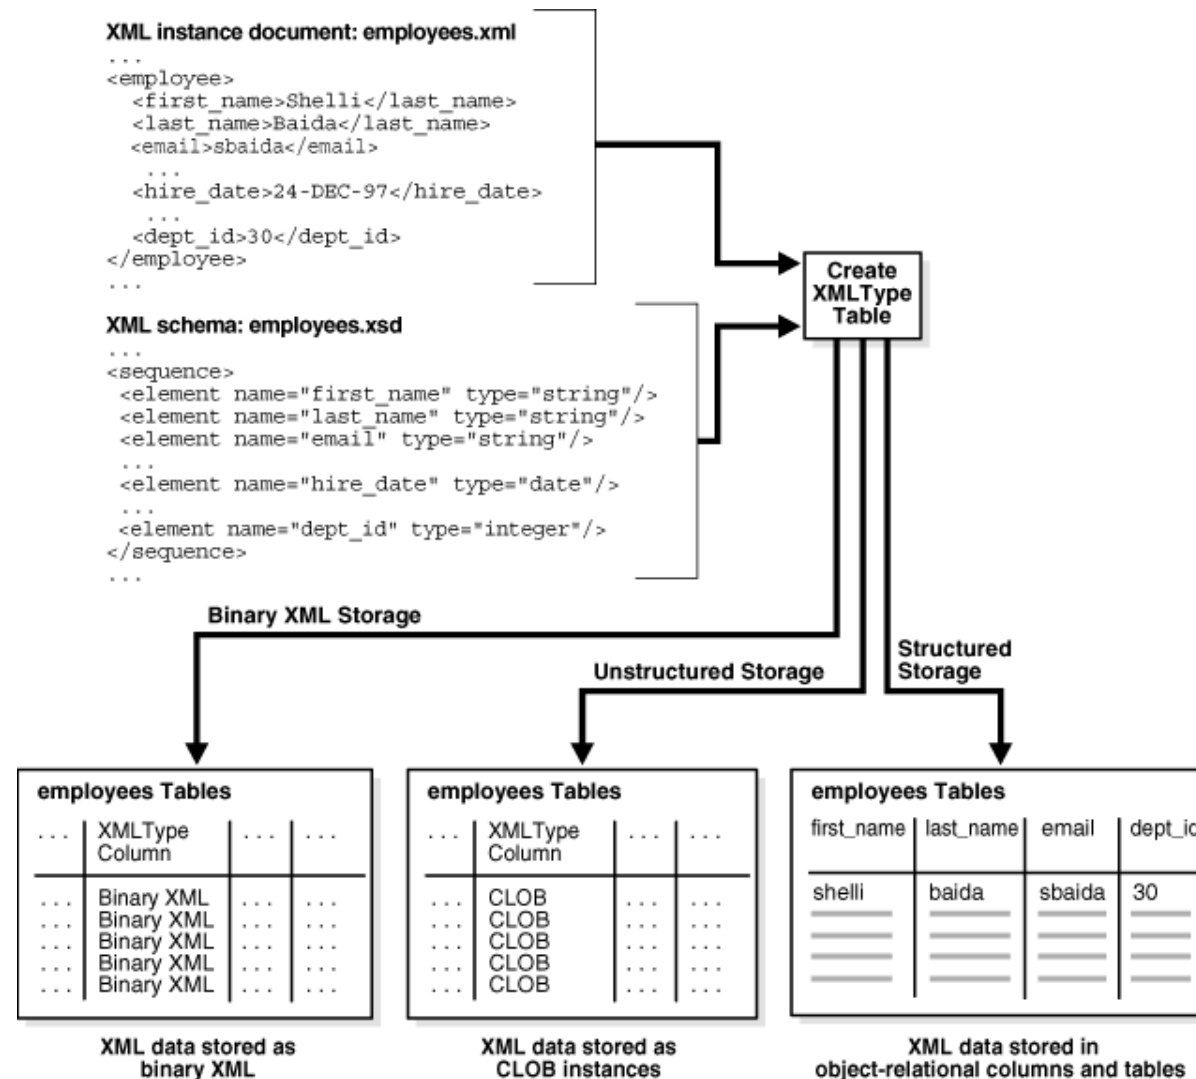
# SQL/XML example

```
CREATE TABLE Customers(  
    CustomerID int PRIMARY KEY,  
    CustomerName nvarchar(100),  
    PurchaseOrders XML, ...  
)
```

# SQL/XML example

```
SELECT CustomerName,  
  
       query(PurchaseOrders,  
  
            'for $p in /po:purchase-order  
  
            where $p/@date < xs:date("2002-10-31")  
  
            return <purchaseorder date="{ $p/@date }">  
  
                { $p/* }  
  
                </purchaseorder>' )  
  
FROM   Customers  
  
WHERE  CustomerID = 42
```

# XML Column Type : 3 Possible Storages



# Oracle XML CLOB Storage

Simplest approach to implement and support

Byte fidelity : preserves original doc (even white space)

Performances

↗ load insert full retrieval

↘ query schema evolution

Need to parse the document for all XML processing  
(memory overhead)

# Oracle XML Relational Storage (OR)

Schema-aware mapping to relational tables

Byte fidelity not always guaranteed

Performances

- ↗ query for highly-structured data
- ↘ query for un-structured data (many joins), full-retrieval
- ↘ flexibility schema evolution, load / insert / delete



# Motivation/Goals for Binary XML

## XML Schema usage


- Need to be efficient for query processing on schemaless & loosely structured schemas; use schema for optimization

## Provide good performance for a wide range of operations

- Query
- DML: Insert/Load, Partial (piecewise) update
- Full-document & fragment retrieval
- Schema Validation & Evolution

# What is binary XML ?

XML tree = balanced parenthesized expression





db book title DMS ) ) author Ramak ) ) author Gehrke ) ) ) )

# What is binary XML ?

## Navigation = find corresponding (open-)closed parenthesis

Update = insert / delete a substring (easy for Oracle's file system)



# What is binary XML ?

BinaryXML is made of sections (subtrees)



Each section has an header recording (among others)

- path leading to that subtree
- sibling-position

<b>Header</b> Path=db.book.title Position = 1	OPEN- ELEMENT	<b>Header</b> Path=db.book.title.text() Position = 1	OPEN- TEXT	DMS	CLOSE- TEXT	CLOSE- ELEMENT
---	------------------	--	---------------	-----	----------------	-------------------



# What is binary XML ?

- Query evaluation : can be done by scanning the headers
- Storage:
  - Compression of elements of the same type (e.g. many authors)
  - Further compression when schema is available
  - Also datatypes (int, string) can have best low-level storage
- Of course, every information (like paths) is encoded with an ID !

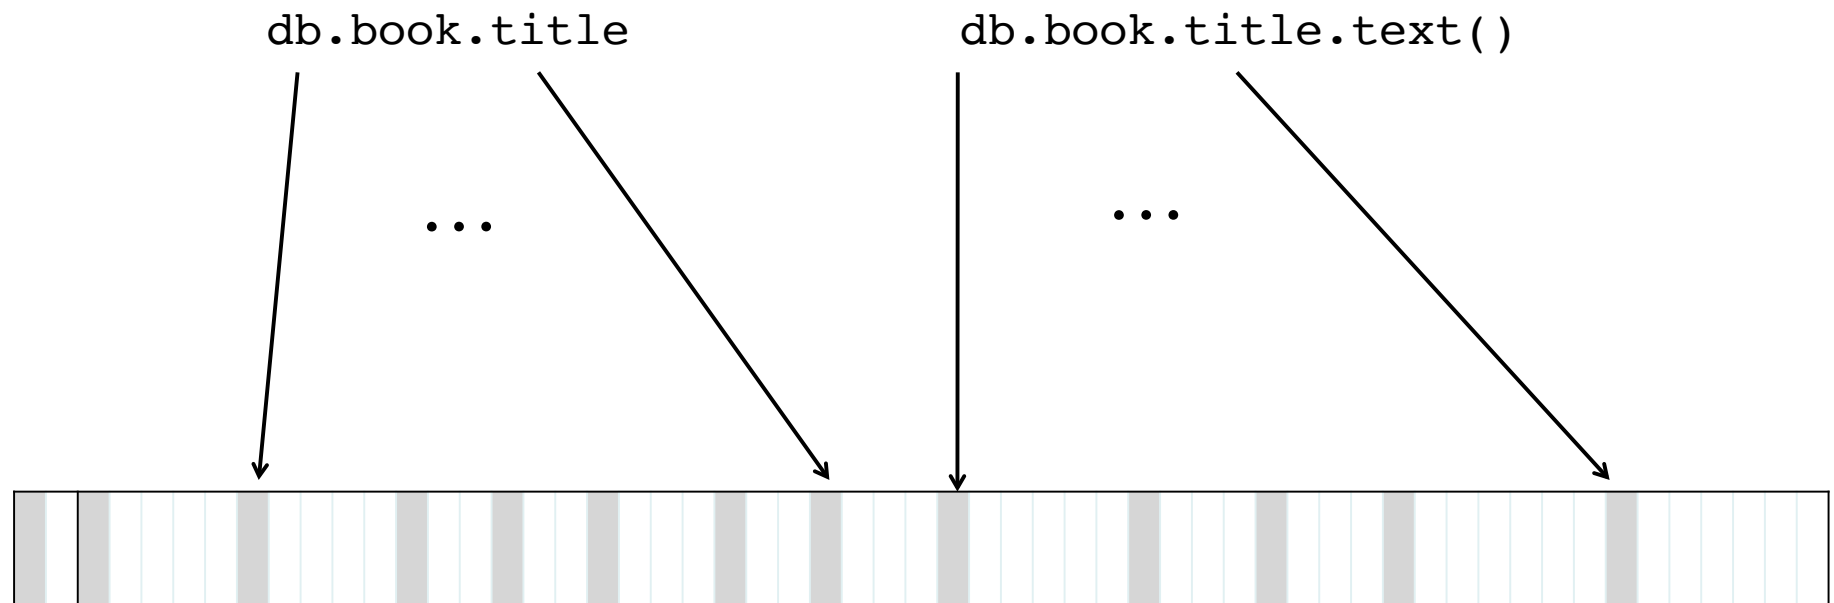
<b>Header</b> Path=db.book.title Position = 1	OPEN- ELEMENT	<b>Header</b> Path=db.book.title.text() Position = 1	OPEN- TEXT	DMS	CLOSE- TEXT	CLOSE- ELEMENT
---	------------------	--	---------------	-----	----------------	-------------------



# What is binary XML ?

Add index to directly access the sections (in constant time).

Dramatically improves performances (at the price of space).




# Comparisons of storage models

	CLOB	OR	Binary XML
Query	poor	excellent	good/excellent
DML	poor	good/excellent	excellent
document retrieval	excellent	good/excellent	excellent
schema flexibility	good	poor	excellent
document fidelity	excellent	poor	good/excellent
mid-tier integration	poor	poor	excellent

# Performance: Compression

compressed  
document size



D1 – Structured

D2 – Semi-structured


D3 – Document-centric

Based on actual ORACLE  
customer datasets

Mix of XML document sizes



# ORACLE Customer use cases



	Data-Centric	Document-Centric	
Use Case	XML schema-based data, with little variation and little structural change over time	Variable, free-form data, with some fixed embedded structures	Variable, free-form data
Typical Data	Employee record	Technical article, with author, date, and title fields	Web document or book chapter
Storage Model	Object-Relational (Structured)	Binary XML	
Indexing	B-tree index	<ul style="list-style-type: none"> <li>- <b>XMLIndex</b> index with structured and unstructured components</li> <li>- XML Full-Text index</li> </ul>	<ul style="list-style-type: none"> <li>- <b>XMLIndex</b> index with unstructured component</li> <li>- XML Full-Text index</li> </ul>

FIGURE 1: XML USE CASES AND XMLTYPE STORAGE MODELS